

Assignment 2: Dask vs PySpark vs Koalas vs Modin *

Emanuel Tomé [†], João Ferreira [‡], Ricardo Faria [§], Vânia Guimarães [¶]

DCC-FCUP, July 2021

1 Introduction

Several parallel DataFrame systems alternatives to Pandas DataFrame are available nowadays. However, the differences between the available libraries is not always clear in which situations one is preferable over the others. Moreover, the cost of transition from Pandas to another library can be different not only for the need (or not) of learning a new syntax, but also in the compatibility of those libraries with others largely used by Data Scientists such as Scikit-Learn and Numpy. In this context, the main objectives of this work are:

- Identify performance bottlenecks when using a specific library.
- Identify syntactic differences among the different libraries.
- Identify operations that are best suited to one particular library.
- Get acquainted with these libraries and knowing what is supported from Pandas, Scikit-learn and Numpy.

2 Brief background on PySpark, Dask, Modin, JobLib, Rapids and Koalas

2.1 Dask

Libraries such as NumPy, Pandas and Scikit-learn, are designed to run on a single core. Therefore, all the data will be temporarily loaded onto the RAM of our local system. However, if we deal with extreme large datasets, we will inevitably problems to run it. Dask appeared mainly to solve this problem. Dask [1] is a parallel computing library that works by distributing larger computations and breaking it down into smaller computations through a task scheduler and task workers. Designed to parallelize in python ecosystems, Dask is suited to solve a wide variety of problems including structured data analysis, large-scale simulations used in scientific computing and general-purpose [2]. However, for data that fits into RAM, Pandas can often be faster and easier to use than Dask DataFrame. The popularity of the Dask is due to the union of the power of distributed computing for data science with the good integration to common Python data tools. Another advantage is that without change the interface, the users can run on clusters with multiple cores or on a common machine using a single process. Otherwise, the users don't really need to worry about the low-level internals, Dask provides several collections for wrapping low-level tasks into high-level workflows.

2.2 Modin

Modin [3] uses Ray or Dask to provide an effortless way to speed up pandas notebooks, scripts and libraries. The authors claim that Modin is able to scale the pandas workflow by changing a single line of code. That is, to use Modin and take advantage of its speedup, the user only has to change the line of code which imports pandas, `import pandas as pd`, to `import modin.pandas as pd` and continues using their previous pandas notebooks.

*This work was submitted on the framework of the course Big Data and Cloud Computing of the Master in Data Science.

[†]Emanuel Tomé is a student at the Faculty of Sciences of the University of Porto. Currently, he is enrolled in the 1st year of the Master's in Data Science (e-mail: up200702634@edu.fc.up.pt).

[‡]João Ferreira is a student at the Faculty of Sciences of the University of Porto. Currently, he is enrolled in the 1st year of the Master's in Data Science (e-mail: up202004115@edu.fc.up.pt)

[§]Ricardo Faria is a student at the Faculty of Sciences of the University of Porto. Currently, he is enrolled in the 1st year of the Master's in Data Science (e-mail: up202004105@edu.fc.up.pt)

[¶]Vânia Guimarães is a student at the Faculty of Sciences of the University of Porto. Currently, she is enrolled in the 1st year of the Master's in Data Science (e-mail: up200505287@edu.fc.up.pt).

Moreover, to use Modin the user does not have to know how many cores their system has and does not need to specify how to distribute the data. Data Scientists spend their time extracting value from their data than on tools that extract data.

2.3 JobLib

Joblib provides solution for several activities, such as loading up large Numpy arrays, persisting python object or performance of custom python function, with the help of parallel computing, memoization and caching mechanism. Joblib [4] provides a set of tools for making the pipeline lightweight to a great extent in Python [5], without dependency on other libraries. Joblib allows to use cache which avoids recomputation of some of the steps and execute parallelization to fully utilise all the cores of CPU/GPU. So, function called with same argument will not be re-compute, instead, output loads back from cache using memmapping. It also provides a compressor during persistence for large data, to save space on disk. A fantastic library that became popular because of its optimized time-complexity feature.

2.4 Rapids

Rapids [6] is a suite of open source software libraries and APIs that provides the ability to execute end-to-end data science and analytics pipelines entirely on GPUs. Rapids utilises NVidia Cuda primitives for low-level compute optimisation, and exposes GPU parallelism and high-bandwidth memory speed through user-friendly Python interfaces. It also focuses on common data preparation tasks for analytics and data science, and includes support for multi-node, multi-GPU deployments, enabling vastly accelerated processing and training on much larger dataset sizes.

2.5 PySpark

Apache Spark is a unified computing engine and a set of libraries for parallel data processing on computer clusters [7]. Spark supports multiple widely used programming languages, such as Python, Java, Scala and R, includes libraries for diverse tasks ranging from SQL to streaming and machine learning and runs anywhere (from a laptop to a cluster of thousands of servers).

PySpark [8] is an interface for Apache Spark in Python. It allows to write Spark applications using Python APIs and also provides the PySpark shell for interactively analyse data in a distributed environment. PySpark supports most of Spark's features such as Spark SQL, DataFrame, Streaming, MLlib (Machine Learning) and Spark Core.

2.6 Koalas

Pandas syntax and PySpark syntax differ to a considerable degree because PySpark has been notably influenced by SQL syntax. Regular pandas users will argue that it is much less intuitive. This is where Koalas enters the picture. Koalas is a data science library that implements the Pandas APIs on top of Apache Spark so, data scientists can use their favourite APIs on datasets of all sizes [9], and be more productive. Pandas does not scale well to big data since it was designed for small data sets. Using Koalas it is possible make the transition from a single machine to a distributed environment without needing to learn a new framework.

3 Materials and methods

3.1 Machines used and their characteristics

As required, we started by creating a cluster similar to i3.4xlarge AWS and a machine similar to i3.16xlarge, on Databricks and Dataproc. However, trial accounts for both platforms do not allow work with such power machines. This information is not provided to us, therefore after we started the experiments, both counts were cancelled. To solve the problem, we asked for a machine at DCC. With advanced project, the server stopped to work. The next alternative was create a single node on Dataproc with 8 CPUs and 32 Gb of memory with 300 dollars, the coupon given by Google for each new member. Although, there was some incompatibilities with Dask and Koalas, that made it impossible to continue the execution of the assignment.

The finally alternative was to create two Virtual Machines. The experiments which required Pandas, Dask, Modin, Koalas and Joblib were executed in a Virtual Machine of Google Cloud Platform with 16 CPUs (Intel Cascade Lake) and 64 GB of memory. Selected zone was us-central1-a.

Another Virtual Machine was created to run the experiments with Rapids. This machine contains 8 CPUs, 30 GB of memory and 1 GPU Nvidia Tesla T4. Selected zone was us-central1-b.

3.2 Datasets description

As we mentioned before, we started with all dataset required. On each new machine created, the data was decreasing, according to the power of the respective machine. At the end, both experiments and the Machine Learning task were executed with 2009, January of Yellow Taxi Trip Records from NYC Taxi and Limousine Commission (TLC) Trip Record Data.

4 Experiment #1: repeat NYC taxi driver dataset study

In this section we present the results obtained for one month of the NYC taxi driver dataset study [9], namely for data from January 2009. Note that only a month of data was considered due to limitations of the available computational capacity. We present the computation times using the following libraries: Pandas (sequential), Joblib, Dask, Modin (with Dask), Koalas (PySpark) and Rapids.

In Figures 1 and 2 are presented the obtained computation times of the standard operations. The first thing to notice is that Rapids has an outstanding performance for all the operations considered. With the exception of operation "count index length" where Pandas is the faster (note that all data fit in memory), Rapids is the faster library. In the opposite side, Joblib seems to be generally the library with worst performance, having poor performance in the majority of operations and the highest total execution time (see Figure 4). However, we should remark that the parallelization with Joblib was made using Pandas dataframes, which should not be the best choice for this library since the Pandas dataframes are stored in columns and not in rows, as for example the Numpy arrays. However, this choice was made since our objective was compare the performance of different libraries when using dataframes.

It is also possible to notice that Modin, followed by Koalas, has by far the worst computation time for the operations "complex arithmetic ops" and "mean of complex arithmetic ops". Indeed, Modin has the second worst total execution time of all benchmarks (see Figure 4).

Considering the geometric mean of the execution times of the standard operations (Figure 4), one can state that Joblib had the worst geometric mean (2.92s) and Rapids and Pandas the best ones. The low computation times for Pandas is only possible since the dataframe fits in memory. We can also notice that although Dask as a lower total computation time than Koalas (Figure 3), the obtained geometric mean for both libraries is similar.

It should be remarked that care should be take when comparing the different computation times since they could not be completely fare or even comparable. For instance, Dask has the lower computation time for reading data. However, in Dask the reading is lazy, which means that the computation time for reading files will be the same regardless the size of the files.

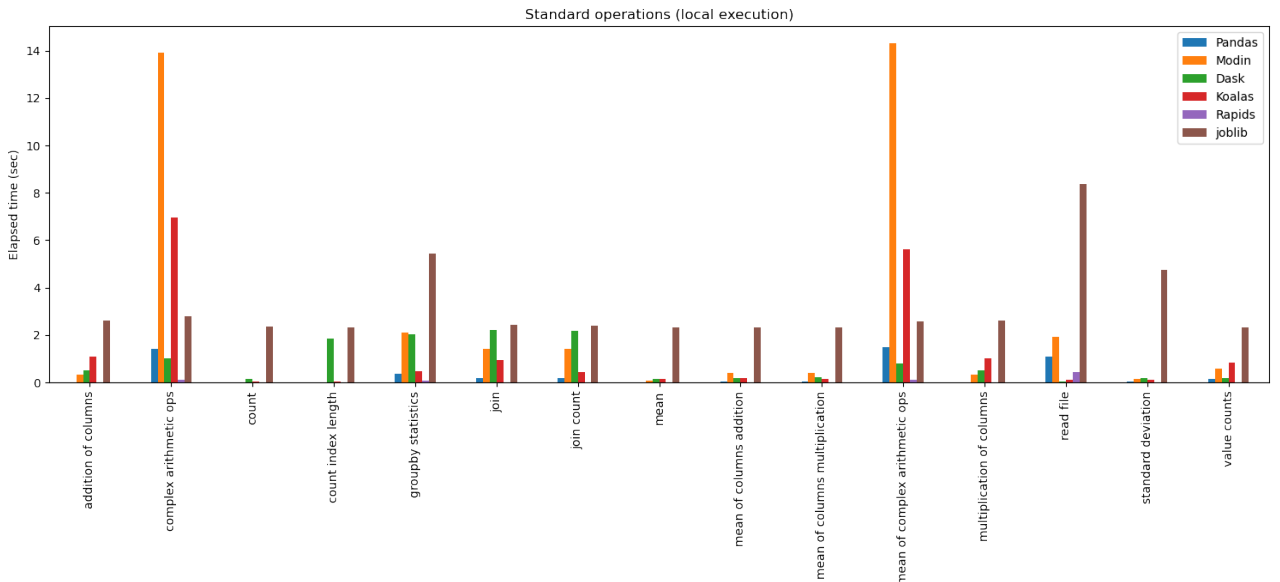


Figure 1: Computation times of the standard operations.

Considering now the computation times obtained for the filtered dataset, that is, for a small dataset, the obtained results are presented in Figures 5 and 6. One can noticed that now Dask has the worst computation times for almost all operations. This confirms what is stated in the documentation of Dask, where the developers confirm that for small datasets does not worth using Dask. Rapids for this smaller dataset is still the faster




Figure 2: Computation times of the standard operations - logarithmic scale.




Figure 3: Total execution time of the standard operations

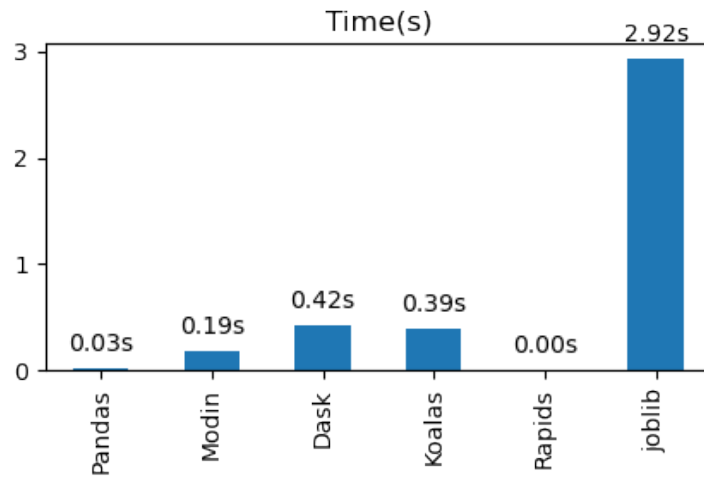


Figure 4: Geometric mean of the execution times of the standard operations.

library closely followed by Pandas. Modin and Koalas have similar total computation times, while Joblib has slightly lower total computation time (see Figure 7).

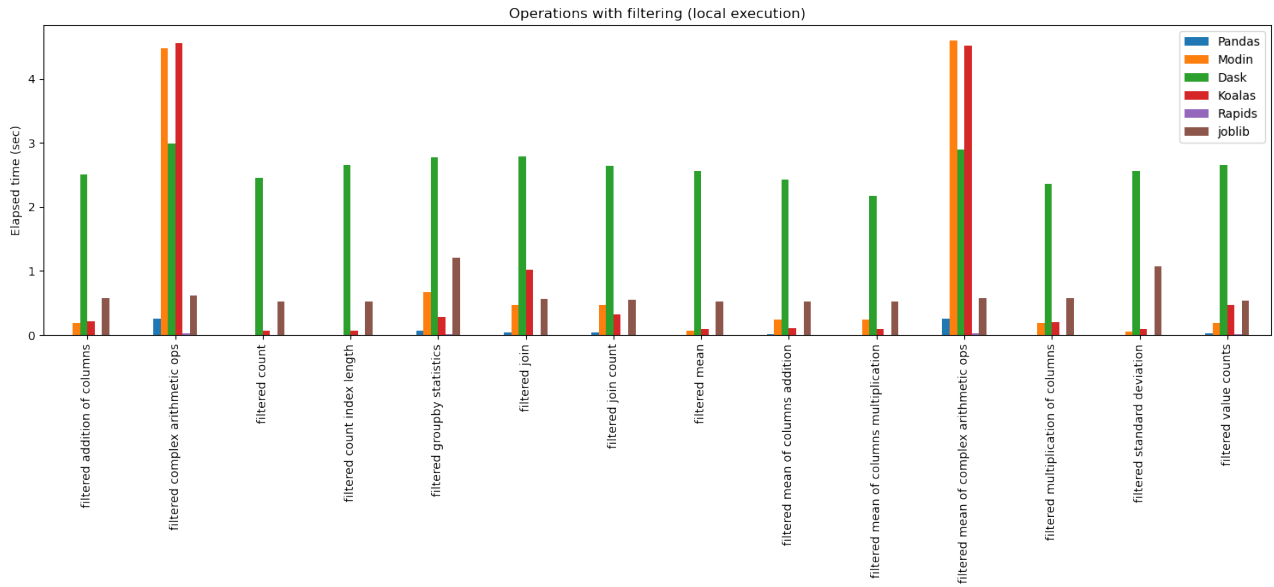


Figure 5: Computation times of the standard operations with filtering.




Figure 6: Computation times of the standard operations with filtering - logarithmic scale.

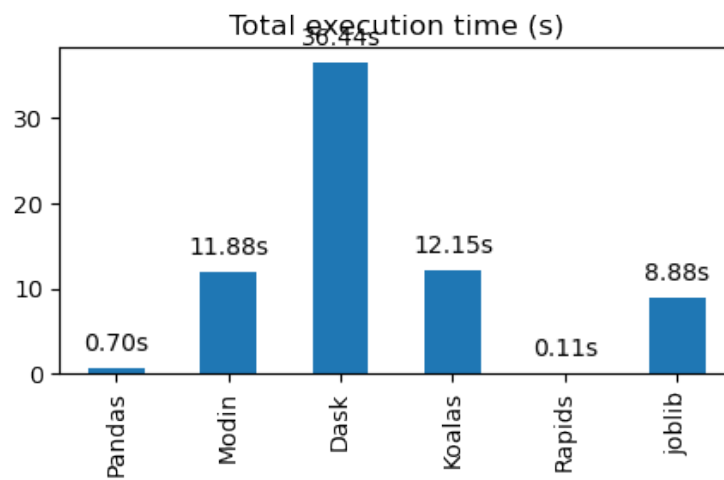


Figure 7: Total execution time of the filtered standard operations.

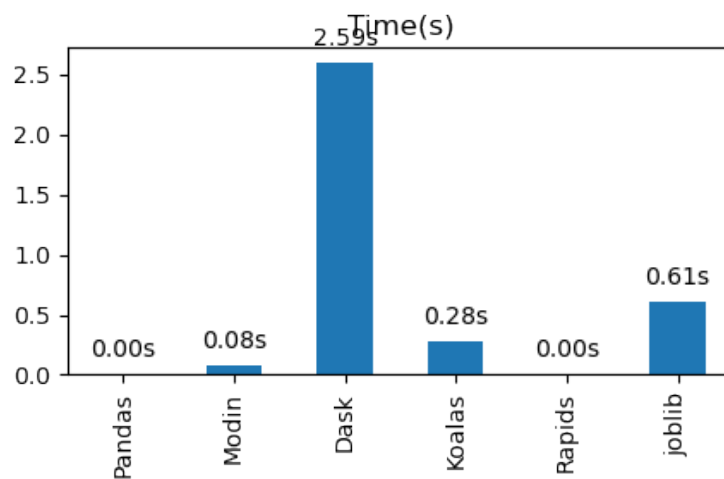


Figure 8: Geometric mean of the execution times of the filtered standard operations.

5 Experiment #2

5.1 Description of the experiment

The aim of Experiment #2 is to build a full machine learning pipeline and predict the target variable "fare_amount" using **Dask**, **Joblib**, **Koalas** and **Rapids**.

A full machine learning pipeline has the following steps:

- Reading the data
- Pre-processing
- Training
- Testing

5.1.1 Read Data

Due to problems of budget and computing resources, the dataset used in the pipeline corresponds to only one month of the NYC taxi driver dataset. The chosen month was January of 2009. With exception of Joblib, which reads the data stored as a Pickle file, all the other libraries read files in the parquet format.

5.1.2 Pre-processing

Depending on the type of dataframe used, the pre-processing may vary but these steps are common to all libraries. Using the time in "trip_pickup_datetime" it was added two new columns: "dayofweek" (0 to 6 according to the day of the week) and "hour". A new variable, "weekday" was created with dayofweek (0 if workday, 1 if weekend). Variable "trip_duration" was created by subtracting "trip_pickup_datetime" from "trip_dropoff_datetime". In "payment_type" some categories had different values with the same mean so these cases were recoded. All values with "CASH" recoded to "Cash" and "CREDIT" recoded to "Credit". Lines with abnormal values were removed according to this selection:

- "fare_amt" > 0
- "trip_distance" > 0
- "trip_duration" > 0
- "tip_amt" ≥ 0

An One Hot Encoding was applied to the categorical variables 'vendor_name' and 'payment_type'. Several variables with no interest were removed from the dataset, some of them are not related with the target, and others had no information. These variables are: 'index', 'trip_pickup_datetime', 'trip_dropoff_datetime', 'rate_code', 'mta_tax', 'store_and_forward', 'vendor_name' and 'payment_type'. The 'total_amt' was also removed since the 'fare_amt' can be easily computed just by subtracting in the 'total_amt' the variables 'tolls_amt' and 'tip_amt'. Finally, it was applied a MinMaxScaler to all numeric variables including the target variable.

5.1.3 Training

After preprocessing, randomSplit was applied to the dataset:

- 70% Train
- 30% Test

The model used by most libraries was XGBRegressor (Extreme Gradient Boosting) of Scikit-learn. We only setup parameter "objective='reg:squarederror'", which is the loss function that is minimized during the training of the model.

XGBRegression is a decision-tree-based ensemble Machine Learning algorithm that applies the principle of boosting. A decision tree is the weak learner, the resulting algorithm is called gradient boosted trees. The "eXtreme" refers to speed enhancements obtained by parallel computation.

For the experiments with Koalas, the Extreme Gradient Boosting was also tested^[10] with Koalas but there were some errors that could not be solved. As described in PySpark Example, the spark-xgboost was installed and imported. However, the function XGBRegressor could not be loaded. Further, inspecting the package installed^[11], several Python versions up to 3.6 are supported. Although, it was not possible to work with XGBRegressor API even using Python 3.5 and 3.6.

Instead of XGB, the model used was GBRegressor (Gradient-Boosted Trees), which is worse in terms of computational time, since it is not optimised with parallel processing. The loss function to be minimised is the same as the XGB models.

5.2 Obtained metrics

Because this is a regression problem, in order to rank the performance of the models that were tested above, in the different pipelines, the metrics used were:

- Mean squared error - MSE.
- Root mean squared error - RMSE.
- Mean absolute error - MAE.
- Max error - ME.
- Coefficient of determination - R^2 .

The MSE was the metric used to fit the models. The results obtained are shown in Table 1.

Table 1: Performance metrics of the fitted models on the different libraries.

	MSE	RMSE	MAE	Max error	R^2	Pipeline execution time (s)
<i>Rapids</i>	0.000244	0.015611	0.013758	0.954854	0.799853	97.75
<i>Koalas</i>	0.000089	0.009446	0.003511	0.989797	0.926805	461.52
<i>Dask</i>	0.000075	0.008655	0.002869	0.992073	0.938329	385.35
<i>Joblib</i>	0.000028	0.005292	0.001278	0.792122	0.977025	3422.39

As we can visualise, the library that obtained the best results (even if the differences are almost insignificant), was the *Joblib*, having the best results among the other ones. The *rapids* library was the one with the worse results and one possible reason for this was a limitation of the *cuML* library, which doesn't currently have an *XGBRegressor* available implemented. To circumvent this, the one used was *XGBRegressor* from *Python XGBoost* package, applied on the dataframe that was saved in GPU.

One additional limitation was the holdout performed on the GPU dataframe, since there isn't currently available a `test_train_split` function to perform the split of the dataframe on the GPU, so the way this was surpassed was by using the indices of the dataframe corresponding to x% of the whole data, without shuffling the data. As a consequence, the method decreased the independence of the data, and increased the bias. Regarding *Koalas*, and as already explained before, since we couldn't use the *XGBRegressor* model with that library, the *GBRegressor* was used. Therefore, care should be taken when comparing the obtained metrics since the models are not exactly the same over the different libraries.

Overall, all the models had a similar performance independently of the library used, however, the execution time of the pipeline is quite different along the libraris. We can see that the *Joblib* was by far the one that had the longest execution time of them all, counting a total of 3422 seconds (57 minutes). Then, we had the *Koalas* library, that counted a total of 462 seconds (approximately 8 minutes), followed by the *Dask* library, which had took about 368 seconds (6 minutes, approximately), and finally, *rapids* which took only 85 seconds (1 minute and 25 seconds). This comes to reinforce the fact that using GPU for machine learning tasks that deal with large datasets is highly efficient, lowering the execution times of almost every standard operation we could think of.

5.3 Profiling

Regarding the profiling of the pipelines, for the majority of the libraries the pipeline part that takes more time to compute is the training of the machine learning model. The exception was *Koalas*, where more computation time was spent in predicting and computing the evaluation metrics. *Joblib* was by far the library that took more time to conclude the computation time. However, it should be noted that this is not directly related with *Joblib*, but with the library/model used to build the model, even though the used model and library trains the model in parallel. If only the pre-processing and "prediction and testing" steps of the pipeline are considered, *Joblib* has a very good performance.

With exception of *Koalas*, where the second bottleneck was the training of the machine learning model, the second bottleneck for all libraries was pre-processing of the data. Indeed, this is a step in the pipeline where several computations take place, such as transformation of the variables, scaling, drop of variables, computation of new variables from others, merging of dataframes, etc. For this reason, it was expected that this stage of the pipeline was one of the stages with higher computation time. The framework with higher computation time was *Dask*. One possible reason for that is the need to create indices in the dataframe of numeric, categorical and target variables, in order to guarantee the correct concatenation of the three. Operations like "set_index" and 'merge/join' are harder to do in a parallel or distributed setting than if they are in-memory on a single machine. Shuffling operations that rearrange data become much more communication intensive.

The operation related to import data, Dask was faster than Rapids that use a GPU. This is due to the fact that one of the features of Dask is lazy execution. That means it loads and then processes the data in chunks, so that only a subset of the data needs to be in memory at any given time. However, cuDF stores Dataframes in GPU memory and uses the GPU to perform computations. Regarding Joblib, this was the library with the worst performance. However, it should be stated that for this library the data was read from a pickle file because the library documentation refers that this is the more efficient way to load data to joblib. If the dataframe was read using, for instance, a lower computation time for this step would be obtained.

Table 2: Line profiling on each of the pipeline’s operations.

	Import data	Pre-processing	Training	Predicting and Testing	Pipeline
<i>Rapids</i>	0.47	2.50	94.16	0.63	97.75
<i>Koalas</i>	2.7	19.04	391.50	443.08	856.34
<i>Dask</i>	0.02	23.92	348.12	13.27	385.35
<i>Joblib</i>	9.51	18.07	3388.54	6.26	3422.39

6 Main difficulties and challenges

In this work we faced many difficulties and challenges. The main ones are listed and described bellow:

- **Rapids limitations** - The lack of some vital functions on the cuDF and cuML libraries made some tasks not as easy as they should. For example, the `train_test_split` function was not implemented and it was needed to implement some kind of algorithm that would split the data into different parts. The problem here was to maintain efficiency, which was not possible due to the sequential programming used to perform this task.
- **XGBRegressor with Koalas** - as already referred, although all the attempts, it was not possible to use the XGBRegressor with Koalas/PySpark. As an alternative, we used the GBRegressor algorithm.
- **Configuration of the Clusters** - In the begging of this assignment we tried to configure a cluster in the Google Cloud Platform (GCP) using Databricks. However, our account was cancelled two times due to the limitation of GCP described in the next point.
- **Limitation in the GCP of using only a CPU with no more than 8 cores** - when using the educational vouchers, it is not possible to use virtual machines with more than 8 cores. If the user does that, the billing and voucher are cancelled.
- **Necessity of changing platform several times** - We start by trying to do our experiments with Databricks. However, in the beginning our account was cancelled and countless options were tried. We then moved to a DCC server that was made available to the authors. However, when we started to implement our experiments related with the machine learning part, the kernel was always stopping. We then moved to a single node in the Dataproc of the GCP where we were not able to run all the libraries. Namely, we where not able to read the parquet files that was created before using koalas. We then finally configured two virtual machines in GCP and these was our final solution. Note that each platform has their own characteristics which requires configuring the system, libraries dependencies and sometimes the code.
- **Scarce resources for a so big dataset** - Besides we have to find ways to load data related to 4 years of records and convert them to parquet format, deal with a system that goes down frequently or fit the data to our current resources without lose the purpose of the project was a big challenge. Additionally, the lack of credits in GCP made it difficult to setup a proper machine to run the tasks.
- **Impossibility of doing distributed execution** - As mentioned on the guidelines of the assignment, we tried to implement a distributed execution on Databricks, tirelessly looking for information that would guide us in this direction as we couldn’t associate the cluster with the Dask client. After a day of intense search, we were informed that it is not possible to do it and that it was no longer necessary to apply the distributed execution in our experiments.
- **Interpretation of the results obtained using the Yappi profiler** - The output from the Yappi profiler is not as simple as the line profiler. While the line profiler outputs the time spent in each command, the yappi profiler outputs the time spent in each operation which makes it hard to identify, without knowing the meaning of each operation, where the program is spending more time. Although yappi may be harder to analyse, it is possible to profile each thread in details.

- **Time limitation** - Despite the extension of the deadline for the assignment, the difficulties encountered along the way were mainly due to the infrastructure needed to execute the experiments. This situation took our time in the analysis of the core of the project.

7 Discussion and Conclusions

In this work we applied different libraries that allow the parallelization of operations in DataFrames and data science pipelines. The obtained results in the benchmarks showed that the use of Rapids/GPUs outperforms by far the results using the other libraries. Dask and Koalas had similar performances and Joblib was the one with the worst computation times. However, it should be state that we used Joblib with dataframes, while it is more suited to use Numpy arrays, which are stored in memory in rows, while the dataframes are stored in columns. Moreover, the kind of operations done in our experiments may not be the ones for which this library is more suited for.

Rapids had an outstanding performance since it used an Nvidia GPU to perform all of the dataframe's operations as well all the machine learning tasks, having some limitations as previously said. Overall, the use of Rapids had an enormous impact in the execution times, decreasing it in relation to the other libraries tested. This was possible because the used GPU made possible to process multiple computations simultaneously, as it has a large number of cores, which allows for better computation of multiple parallel processes. This can be more specifically explained by the fact that GPUs are bandwidth optimised while CPUs are latency optimised, i.e the GPU can fetch much more memory at once, while CPU's can only fetch small amounts of memory quickly, having to perform a lot of memory operations to achieve the same goal. When you add this up to thread parallelism, GPUs get a lot of advantage from the fact that now it can perform large amounts of data operations, simultaneously and multiple times.

Koalas has the disadvantage to be the library with the more distinct syntax of the ones tested, specially when it is necessary to work in PySpark dataframes. Indeed, the Koalas syntax is similar to pandas syntax. However, there was the need of working directly with PySpark since there are functions that are not implemented in Koalas. Even though, since this is a brand new technology this may improve over the time and more functionalities and documentation may be provided.

Regarding the line profiling, we concluded that, at least for the model used, for the majority of the libraries the biggest part of the pipeline computation time is spent training the model. The exception was the Koalas library, which has more computation time in prediction and testing. The prediction and testing step was computed again, but this time only for predicting and it's computation time was 0.38s. The prediction time is low but the time used to compute RegressionEvaluator for each metric is quite large.

Finally, it should be referred that a bigger dataset may have been used to due a more fair comparison between the different libraries. Indeed, a data set of about 2.5GB may not be enough to evaluate the potential of the studied libraries. Therefore, this study may be complemented in the future with other analysis, operations and a bigger datasets.

References

- [1] Dask Development Team. *Dask: Library for dynamic task scheduling*. 2016. URL: <https://dask.org>.
- [2] J. Daniel. *Data Science with Python and Dask*. Manning Publications, 2019. ISBN: 9781638353546. URL: <https://books.google.pt/books?id=KTkzEAAAQBAJ>.
- [3] Devin Petersohn et al. "Towards Scalable Dataframe Systems". In: *CoRR* abs/2001.00888 (2020). arXiv: 2001.00888. URL: <http://arxiv.org/abs/2001.00888>.
- [4] Joblib Development Team. *Joblib: running Python functions as pipeline jobs*. 2020. URL: <https://joblib.readthedocs.io/>.
- [5] Pratik Gandhi. *Using joblib to speed up your Python pipelines*. Sept. 2020. URL: <https://towardsdatascience.com/using-joblib-to-speed-up-your-python-pipelines-dd97440c653d>.
- [6] Rapids Development Team. *About Rapids*. 2021. URL: <https://rapids.ai/about.html>.
- [7] B. Chambers and M. Zaharia. *Spark: The Definitive Guide: Big Data Processing Made Simple*. O'Reilly Media, 2018. ISBN: 9781491912300. URL: <https://books.google.pt/books?id=oitLDwAAQBAJ>.
- [8] *PySpark Documentation*. URL: <https://spark.apache.org/docs/latest/api/python/index.html>.
- [9] *How fast Koalas and PySpark are compared to Dask - The Databricks Blog*. Apr. 2021. URL: <https://databricks.com/blog/2021/04/07/benchmark-koalas-pyspark-and-dask.html>.
- [10] *Databricks XGBoost Example*. URL: <https://databricks.com/notebooks/xgboost/xgboost4j-spark-example.html>.

[11] *Github Spark-xgboost*. URL: <https://github.com/sllynn/spark-xgboost/blob/master/setup.py>.

Appendixes

A Detailed description of the used Machines

Machine 1 - No GPU

CPU:

- Architecture: x86_64
- CPU op-mode(s): 32-bit, 64-bit
- Byte Order: Little Endian
- CPU(s): 16
- On-line CPU(s) list: 0-15
- Thread(s) per core: 2
- Core(s) per socket: 8
- Socket(s): 1
- NUMA node(s): 1
- Vendor ID: GenuineIntel
- CPU family: 6
- Model: 85
- Model name: Intel(R) Xeon(R) CPU
- Stepping: 7
- CPU MHz: 3100.240
- BogomIPS: 6200.48
- Hypervisor vendor: KVM
- Virtualization type: full
- L1d cache: 32K
- L1i cache: 32K
- L2 cache: 1024K
- L3 cache: 25344K
- NUMA node0 CPU(s): 0-15

RAM: 64GB

Machine 2 - with GPU

- Architecture: x86_64
- CPU op-mode(s): 32-bit, 64-bit
- Byte Order: Little Endian
- CPU(s): 8
- On-line CPU(s) list: 0-7
- Thread(s) per core: 2

- Core(s) per socket: 4
- Socket(s): 1
- NUMA node(s): 1
- Vendor ID: GenuineIntel
- CPU family: 6
- Model: 63
- Model name: Intel(R) Xeon(R) CPU @ 2.30GHz
- Stepping: 0
- CPU MHz: 2299.998
- BogomIPS: 4599.99
- Hypervisor vendor: KVM
- Virtualization type: full
- L1d cache: 32K
- L1i cache: 32K
- L2 cache: 256K
- L3 cache: 46080K
- NUMA node0 CPU(s): 0-7

RAM: 30GB

GPU: NVIDIA Tesla T4 (16GB)

B Benchmark results

C Source code of the setup functions for JobLib

```

1 import os
2 import joblib
3 from joblib import dump, load
4 from joblib import Parallel, delayed
5 import pandas as pd
6 import numpy as np
7 from tqdm import tqdm
8 import functools
9 # import databricks.koalas as ks
10 import time
11 import math
12
13 def get_results(benchmarks):
14     """Return a pandas DataFrame containing benchmark results."""
15     return pd.DataFrame.from_dict(benchmarks)
16
17 def benchmark(f, n_jobs=-1, df=None, benchmarks=None, name=None, **kwargs):
18     """Benchmark the given function against the given DataFrame.
19
20     Parameters
21     -----
22     f: function to benchmark
23     df: data frame
24     benchmarks: container for benchmark results
25     name: task name
26
27     Returns
28     -----
29     Duration (in seconds) of the given operation
30     """
31     start_time = time.time()

```

Table 3: Computation times of the benchmark operations (in seconds).

	Pandas	Modin	Dask	Koalas	Rapids	Joblib
read file	1.097639	1.933903	0.021552	0.116975	0.440296	8.386324
count	0.000021	0.000017	0.161452	0.048637	0.000008	2.341892
count index length	0.000002	0.000004	1.84509	0.053336	0.000003	2.322545
mean	0.017941	0.090053	0.163021	0.144577	0.00205	2.328575
standard deviation	0.035824	0.136144	0.173782	0.102439	0.001807	4.737141
mean of columns addition	0.029497	0.395677	0.19121	0.172799	0.004861	2.321197
addition of columns	0.011179	0.315962	0.493689	1.09793	0.001438	2.600978
mean of columns multiplication	0.029217	0.384626	0.202049	0.141332	0.004349	2.319647
multiplication of columns	0.011427	0.315781	0.508529	1.015454	0.001162	2.606427
value counts	0.156386	0.576157	0.168916	0.835458	0.016973	2.336677
mean of complex arithmetic ops	1.504419	14.310768	0.784181	5.617018	0.112237	2.557495
complex arithmetic ops	1.431074	13.929066	1.032044	6.947288	0.114976	2.783355
groupby statistics	0.372646	2.109774	2.039868	0.481217	0.082091	5.423327
join count	0.177912	1.429141	2.177256	0.420335	0.015542	2.403003
join	0.177711	1.428062	2.215343	0.932832	0.011545	2.440177
filtered count	0.000005	0.000017	2.459231	0.069556	0.000009	0.520608
filtered count index length	0.000001	0.000004	2.660978	0.069706	0.000004	0.519908
filtered mean	0.003317	0.061901	2.558253	0.095373	0.001124	0.523736
filtered standard deviation	0.006458	0.058423	2.55637	0.09805	0.000816	1.067333
filtered mean of columns addition	0.00735	0.241343	2.422405	0.113718	0.00173	0.525848
filtered addition of columns	0.001515	0.185699	2.506612	0.218986	0.000979	0.577525
filtered mean of columns multiplication	0.004525	0.245141	2.175726	0.100009	0.001773	0.525475
filtered multiplication of columns	0.00138	0.191738	2.362341	0.208008	0.001014	0.576129
filtered mean of complex arithmetic ops	0.252189	4.605596	2.900001	4.521427	0.033652	0.578335
filtered complex arithmetic ops	0.250223	4.477522	2.987706	4.55545	0.032899	0.619596
filtered value counts	0.027867	0.192825	2.652701	0.468087	0.011148	0.532077
filtered groupby statistics	0.073863	0.677037	2.770664	0.283749	0.018422	1.203191
filtered join count	0.03491	0.474413	2.639927	0.328058	0.00389	0.544891
filtered join	0.033922	0.470783	2.7875	1.020233	0.003434	0.561686

```

32     ret = func_parallel(f, n_jobs, df, **kwargs)
33     benchmarks['duration'].append(time.time() - start_time)
34     benchmarks['task'].append(name)
35     print(f"{name} took: {benchmarks['duration'][-1]} seconds")
36
37     return benchmarks['duration'][-1]
38
39
40 def func_parallel(f, n_jobs=-1, df=None, **kwargs):
41     n_jobs = n_jobs if n_jobs>0 else os.cpu_count()
42
43     if df is not None:
44         dim0 = df.shape[0] # numero de linhas
45         chunk = math.ceil(dim0/n_jobs) #if n_jobs else math.ceil(dim0/os.cpu_count()) #
46         tamanho do chunk
47         nb_chunks = math.ceil(dim0/chunk)
48
49         index_list = [(i*chunk,min(chunk*(i+1),dim0)) for i in range(nb_chunks)]
50
51     if f.__name__ in ['standard_deviation']:
52         average = func_parallel(mean, n_jobs=n_jobs, df=df)
53
54         res_list = Parallel(n_jobs=n_jobs)(delayed(f)(df.iloc[index[0]:index[1],:], average, **
55         kwargs) for index in index_list)
56
57         res = (np.sum([res1[0] for res1 in res_list])/(np.sum([res1[1] for res1 in res_list])
58         -1))*(1/2)
59
60     elif f.__name__ in ['write_file_joblib']:
61         Parallel(n_jobs=n_jobs)(delayed(f)(df.iloc[index[0]:index[1],:], i, pickle_file=kwargs
62         ['pickle_file']) for i,index in enumerate(index_list))
63         res = 1
64
65     elif f.__name__ in ['read_file_parquet']:
66         res_list=Parallel(n_jobs=n_jobs)(delayed(f)(pickle_file) for pickle_file in kwargs['
67         files_list'])
68         res = pd.concat(res_list)
69
70     else:
71         res_list = Parallel(n_jobs=n_jobs)(delayed(f)(df.iloc[index[0]:index[1],:], **kwargs)
72         for index in index_list)
73
74         if f.__name__ in ['count', 'count_index_length']:
75             res = sum(res_list)
76
77         elif f.__name__ in ['mean','mean_of_sum','mean_of_product','
78         mean_of_complicated_arithmetic_operation']:
79             res = np.sum([res1[0] for res1 in res_list])/np.sum([res1[1] for res1 in res_list
80             ])
81
82         elif f.__name__ in ['sum_columns','product_columns','complicated_arithmetic_operation'
83         ]:
84             res = pd.concat(res_list)
85
86         elif f.__name__ in ['value_counts']:
87             res = functools.reduce(lambda a, b: a.add(b, fill_value=0), res_list)
88
89         elif f.__name__ in ['groupby_statistics']:
90             # mean computation
91             res = functools.reduce(lambda a, b: a.add(b, fill_value=0), res_list)
92             mean_fare = pd.DataFrame(res.fare_amt['sum']/res.fare_amt['count'], columns=['
93             fare_amt_mean'])
94             mean_tip = pd.DataFrame(res.tip_amt['sum']/res.tip_amt['count'], columns=['
95             tip_amt_mean'])
96
97             # std computation
98             def groupby_statistics_2(df, mean_fare, mean_tip):
99                 gb = df.groupby(by='passenger_count')
100                 df_2 = pd.DataFrame()
101                 for i, data in gb:
102                     df_2.loc[i, 'fare_amt_std'] = sum((data['fare_amt'] - mean_fare.loc[i, '
103                     fare_amt_mean'])**2)
104                     df_2.loc[i, 'tip_amt_std'] = sum((data['tip_amt'] - mean_tip.loc[i, '
105                     tip_amt_mean'])**2)
106                 return df_2

```

```

95         res_list_std = Parallel(n_jobs=n_jobs)(delayed(groupby_statistics_2)(df.iloc[index
96 [0]:index[1],:], mean_fare, mean_tip) for index in index_list)
97         res_std = functools.reduce(lambda a, b: a.add(b, fill_value=0), res_list_std)
98         res_std['fare_amt_std']=(res_std['fare_amt_std']/(res['fare_amt']['count']-1))
99         res_std['tip_amt_std']=(res_std['tip_amt_std']/(res['tip_amt']['count']-1))*(1/2)
100
101         res = pd.concat([mean_fare, res_std['fare_amt_std'], mean_tip, res_std['
102 tip_amt_std']],axis=1)
103         res.index.name= 'passenger_count'
104
105         elif f.__name__ in ['join_data', 'join_count']:
106             res = res_list[0]
107
108     return res

```

D Source code of standard operations for JobLib

```

1 def write_file_joblib(df, i, pickle_file):
2     with open(pickle_file+'_'+str(i), 'wb') as f:
3         dump(df, f)#, compress='zlib')
4
5 def read_file_parquet(pickle_file):
6     with open (pickle_file, 'rb') as f:
7         df = load(f)
8     return df
9
10 def count(df=None): # DONE
11     return len(df)
12
13 def count_index_length(df=None): # DONE
14     return len(df.index)
15
16 def mean(df): # DONE
17     return (df.fare_amt.sum(), df.shape[0])
18
19 def standard_deviation(df, average): # DONE
20     res = (df.fare_amt-average)**2
21     return (res.sum(), res.shape[0])
22
23 def mean_of_sum(df): # DONE
24     return ((df.fare_amt + df.tip_amt).sum(), df.shape[0])
25
26 def sum_columns(df): # DONE
27     x = df.fare_amt + df.tip_amt
28     return x
29
30 def mean_of_product(df): # DONE
31     return ((df.fare_amt * df.tip_amt).sum(), df.shape[0])
32
33 def product_columns(df): # DONE
34     x = df.fare_amt * df.tip_amt
35     return x
36
37 def value_counts(df): # DONE
38     val_counts = df.fare_amt.copy().value_counts()
39     return val_counts
40
41 def complicated_arithmetic_operation(df): # DONE
42     theta_1 = df.start_lon
43     phi_1 = df.start_lat
44     theta_2 = df.end_lon
45     phi_2 = df.end_lat
46     temp = (np.sin((theta_2 - theta_1) / 2 * np.pi / 180) ** 2
47             + np.cos(theta_1 * np.pi / 180) * np.cos(theta_2 * np.pi / 180) * np.sin((phi_2 -
48 phi_1) / 2 * np.pi / 180) ** 2)
49     ret = np.multiply(np.arctan2(np.sqrt(temp), np.sqrt(1-temp)),2)
50     return ret
51
52 def mean_of_complicated_arithmetic_operation(df): # DONE
53     theta_1 = df.start_lon
54     phi_1 = df.start_lat
55     theta_2 = df.end_lon

```

```

55 phi_2 = df.end_lat
56 temp = (np.sin((theta_2 - theta_1) / 2 * np.pi / 180) ** 2
57         + np.cos(theta_1 * np.pi / 180) * np.cos(theta_2 * np.pi / 180) * np.sin((phi_2 -
58         phi_1) / 2 * np.pi / 180) ** 2)
59 ret = np.multiply(np.arctan2(np.sqrt(temp), np.sqrt(1-temp)), 2)
60 return (ret.sum(), ret.shape[0])
61
62 def groupby_statistics(df): # DONE
63     gb = df.groupby(by='passenger_count').agg(
64         {
65             'fare_amt': ['sum', 'count'],
66             'tip_amt': ['sum', 'count']
67         }
68     )
69     return gb
70
71 def join_count(df, **kwargs): # DONE
72     return len(pd.merge(df, kwargs['other'], left_index=True, right_index=True))
73
74 def join_data(df, **kwargs): # DONE
75     return pd.merge(df, kwargs['other'], left_index=True, right_index=True)

```

E Source code of Experiment #2

E.1 Dask

```

1 # Read data
2 import databricks.koalas as ks
3 def read_data():
4
5     t0 = time.time()
6     dask_data = dd.read_parquet('../data/2009-01/', index_col='index')
7     t1 = time.time()
8     print(f'[INFO] Computation time: {t1-t0}s')
9
10    return dask_data
11
12 def pre_processing(df):
13
14     # Create variables dayofweek and hour
15     df=df.assign(dayofweek = df.trip_pickup_datetime.dt.dayofweek,
16                 hour = df.trip_pickup_datetime.dt.hour)
17
18     # Creation of the variable weekend (0-weekend, 1-work day)
19     df['weekend'] = (df['dayofweek'] < 5).astype(int)
20
21     # Creation of the variables trip_duration
22     df['trip_duration'] = (df['trip_dropoff_datetime'] - df['trip_pickup_datetime'])/timedelta
23     (minutes=1)
24     #df['trip_duration'] = df['trip_duration']/timedelta(seconds=1)
25
26     # Drop variables with no interest
27     df = df.drop(['trip_pickup_datetime',
28                 'trip_dropoff_datetime',
29                 'rate_code', 'mta_tax',
30                 'store_and_forward', 'dayofweek', 'total_amt'], axis = 1)
31
32     # Corretion of the classes of the variable payment_type
33     df = df.categorize(columns=['vendor_name', 'payment_type'])
34     df['payment_type'] = df['payment_type'].replace('CASH', 'Cash')
35     df['payment_type'] = df['payment_type'].replace('CREDIT', 'Credit')
36
37
38     # Remove lines with abnormal values
39     df = df[df['fare_amt']>0]
40     df = df[df['trip_distance']>0]
41     df = df[df['trip_duration']>0]
42     df = df[df['tip_amt']>=0]
43
44     # Min-max scaling of numerical variables
45     numeric_variables = df[['passenger_count', 'trip_distance', 'start_lon', 'start_lat', '
46     end_lon',

```



```

47         'trip_duration', 'fare_amt']]
48
49     scaler = MinMaxScaler()
50     scaler.fit(numeric_variables)
51     num_data = scaler.transform(numeric_variables)
52     num_data = num_data.assign(idx=1)
53     num_data = num_data.set_index(num_data.idx.cumsum()-1)
54     num_data = num_data.drop('idx', axis=1)
55
56     # One hot encoding of variables vendor_name and payment_type
57     categorical_variables = df[['vendor_name', 'payment_type']]
58     cat_data = pd.get_dummies(categorical_variables)
59     cat_data = cat_data.assign(idx=1)
60     cat_data = cat_data.set_index(cat_data.idx.cumsum()-1)
61     cat_data = cat_data.drop('idx', axis=1)
62
63     weekend = df[['weekend']]
64     weekend = weekend.assign(idx=1)
65     weekend = weekend.set_index(weekend.idx.cumsum()-1)
66     weekend = weekend.drop('idx', axis=1)
67
68     # Concatenate numerical variables and categorical variables
69     num_cat_data = pd.merge(num_data, cat_data, left_index=True, right_index=True)
70     dask_data = pd.merge(num_cat_data, weekend, on = 'idx')#left_index=True, right_index=True
71     ) #
72
73     return dask_data
74
75 def training(dask_data):
76     # Split into training and testing data
77     train, test = dask_data.random_split([0.7, 0.3], random_state=0)
78
79     # Separate labels from data
80     y_train = train.fare_amt
81     y_test = test.fare_amt
82
83     del train['fare_amt'] # remove informative column from data
84     del test['fare_amt'] # remove informative column from data
85
86     model = XGBRegressor(objective='reg:squarederror')
87
88     t0=time.time()
89
90     model.fit(train, y_train)
91
92     t1 = time.time()
93     print(f'[INFO] Computation time: {t1-t0}')
94
95     return model, train, y_train, test, y_test
96
97 folder_out = '2009-01/'
98
99 def save_model(model, model_name):
100     io.save_obj(model,model_name)
101
102 #save_model(model, folder_out + 'Dask_model'+ datetime.now().strftime("%H%M%S"))
103
104 # Prediction and test
105 from dask_ml.metrics import mean_absolute_error
106
107 def pred_test(model, test, y_test):
108     # Predictions
109     prediction = model.predict(test)
110
111     # Evaluation metrics computation
112     y_test = np.asarray(y_test)
113     prediction = np.asarray(prediction)
114
115     # Metrics
116     df_metrics = pd.DataFrame()
117     df_metrics.loc[0,'max_error'] = metrics.max_error(y_test, prediction)
118     df_metrics.loc[0,'R^2'] = metrics.explained_variance_score(y_test,prediction)
119     df_metrics.loc[0,'MAE'] = metrics.mean_absolute_error(y_test,prediction)
120     df_metrics.loc[0,'MSE'] = metrics.mean_squared_error(y_test,prediction)
121     df_metrics.loc[0,'RMSE'] = np.sqrt(metrics.mean_squared_error(y_test,prediction))

```

```

122     return df_metrics
123
124
125 def pipeline_dask(n_jobs=-1):
126     t0 = time.time()
127     # Read data
128     print('[INFO] Reading data...')
129     df = read_data()
130     print('[INFO] Data loaded.')
131
132     # Pre-processing
133     print('[INFO] Pre-processing')
134     dask_data = pre_processing(df)
135
136     #Train the model
137     print('[INFO] Training the model...')
138     t2 = time.time()
139     model, train, y_train, test, y_test = training(dask_data)
140     t3 = time.time()
141     print(f'[INFO] Model trained in {t3-t2} s.')
142
143     # Save the model
144     print('[INFO] Saving the model...')
145     save_model(model, folder_out + 'Dask_model'+ datetime.now().strftime("%H%M%S"))
146     print('[INFO] Model saved.')
147
148     # Predict and model metrics
149     print('[INFO] Predicting and computing model metrics...')
150     df_metrics = pred_test(model, test, y_test)
151
152     # Save metrics
153     io.save_obj(df_metrics, folder_out + 'Dask_model_metrics'+ datetime.now().strftime("%H%M%S"))
154     t1 = time.time()
155     print('[INFO] All metrics computed.\n[INFO] End of the pipeline.')
156     print(f'[INFO] Computation time for training the model: {t1-t0}s')

```

E.2 Joblib

```

1 def func_parallel(f, n_jobs=-1, df=None, **kwargs):
2     n_jobs = n_jobs if n_jobs>0 else os.cpu_count()
3
4     if df is not None:
5         dim0 = df.shape[0] # numero de linhas
6         chunk = math.ceil(dim0/n_jobs) #if n_jobs else math.ceil(dim0/os.cpu_count()) #
        tamanho do chunk
7         nb_chunks = math.ceil(dim0/chunk)
8
9         index_list = [(i*chunk,min(chunk*(i+1),dim0)) for i in range(nb_chunks)]
10
11
12     if f.__name__ in ['standard_deviation']:
13         average = func_parallel(mean, n_jobs=n_jobs, df=df)
14
15         res_list = Parallel(n_jobs=n_jobs)(delayed(f)(df.iloc[index[0]:index[1],:], average, **
        kwargs) for index in index_list)
16
17         res = (np.sum([res1[0] for res1 in res_list])/(np.sum([res1[1] for res1 in res_list])
        -1))*(1/2)
18
19     elif f.__name__ in ['write_file_joblib']:
20         Parallel(n_jobs=n_jobs)(delayed(f)(df.iloc[index[0]:index[1],:], i, pickle_file=kwargs
        ['pickle_file']) for i,index in enumerate(index_list))
21         res = 1
22
23     elif f.__name__ in ['read_file_parquet']:
24         res_list=Parallel(n_jobs=n_jobs)(delayed(f)(pickle_file) for pickle_file in kwargs['
        files_list'])
25         res = pd.concat(res_list)
26
27     else:
28         res_list = Parallel(n_jobs=n_jobs)(delayed(f)(df.iloc[index[0]:index[1]], **kwargs)
        for index in index_list)
29
30         if f.__name__ in ['count', 'count_index_length']:

```

```

31         res = sum(res_list)
32
33         elif f.__name__ in ['mean', 'mean_of_sum', 'mean_of_product', '
mean_of_complicated_arithmetic_operation']:
34             res = np.sum([res1[0] for res1 in res_list])/np.sum([res1[1] for res1 in res_list
35 ])
36
37         elif f.__name__ in ['sum_columns', 'product_columns', 'complicated_arithmetic_operation'
38 ]:
39             res = pd.concat(res_list)
40
41         elif f.__name__ in ['value_counts']:
42             res = functools.reduce(lambda a, b: a.add(b, fill_value=0), res_list)
43
44         elif f.__name__ in ['groupby_statistics']:
45             # mean computation
46             res = functools.reduce(lambda a, b: a.add(b, fill_value=0), res_list)
47             mean_fare = pd.DataFrame(res.fare_amt['sum']/res.fare_amt['count'], columns=['
48 fare_amt_mean'])
49             mean_tip = pd.DataFrame(res.tip_amt['sum']/res.tip_amt['count'], columns=['
50 tip_amt_mean'])
51
52             # std computation
53             def groupby_statistics_2(df, mean_fare, mean_tip):
54                 gb = df.groupby(by='passenger_count')
55                 df_2 = pd.DataFrame()
56                 for i, data in gb:
57                     df_2.loc[i, 'fare_amt_std'] = sum((data['fare_amt'] - mean_fare.loc[i, '
58 fare_amt_mean'])**2)
59                     df_2.loc[i, 'tip_amt_std'] = sum((data['tip_amt'] - mean_tip.loc[i, '
60 tip_amt_mean'])**2)
61                 return df_2
62
63             res_list_std = Parallel(n_jobs=n_jobs)(delayed(groupby_statistics_2)(df.iloc[index
64 [0]:index[1],:], mean_fare, mean_tip) for index in index_list)
65             res_std = functools.reduce(lambda a, b: a.add(b, fill_value=0), res_list_std)
66             res_std['fare_amt_std']=(res_std['fare_amt_std']/(res['fare_amt']['count']-1))
67             res_std['tip_amt_std']=(res_std['tip_amt_std']/(res['tip_amt']['count']-1))**2)
68
69             res = pd.concat([mean_fare, res_std['fare_amt_std'], mean_tip, res_std['
70 tip_amt_std']],axis=1)
71             res.index.name= 'passenger_count'
72
73         elif f.__name__ in ['join_data', 'join_count']:
74             res = res_list[0]
75
76         return res
77
78 def read_file_parquet(pickle_file):
79     with open (pickle_file, 'rb') as f:
80         df = load(f)
81     return df
82
83 def mean(df): # DONE
84     return (df.sum(), df.shape[0])
85
86 # Read data
87 def read_data():
88     folder_in_out = '2009-01/'
89     folder_data = '../data/joblib_taxi_pickles/' + folder_in_out
90     files_list = ['part_0', 'part_1', 'part_2', 'part_3', 'part_4', 'part_5', 'part_6',
91 'part_7']
92     files_list =[folder_data+file for file in files_list]
93
94     t0 = time.time()
95
96     joblib_data = func_parallel(read_file_parquet, n_jobs=-1, df=None, name='read file',
97 files_list=files_list)
98
99     t1 = time.time()
100     print(f'[INFO] Computation time: {t1-t0}s')
101
102     return joblib_data
103
104 # Pre-processing

```

```

95 from datetime import timedelta
96 from sklearn.preprocessing import MinMaxScaler
97
98 def pre_processing_1(df):
99     # Create variables dayofweek and hour
100     df=df.assign(dayofweek = df.trip_pickup_datetime.dt.dayofweek,
101                 hour = df.trip_pickup_datetime.dt.hour)
102
103     # Creation of the variable weekend (0-weekend, 1-work day)
104     df['weekend'] = (df['dayofweek'] < 5).astype(int)
105
106     # Drop of variable weekend
107     df = df.drop(['dayofweek'], axis=1)
108
109     # Creation of the variables trip_duration
110     df['trip_duration'] = (df['trip_dropoff_datetime'] - df['trip_pickup_datetime'])/timedelta
111     (seconds=1)
112
113     # Remove lines with abnormal values
114     df = df[df['fare_amt']>0]
115     df = df[df['trip_distance']>0]
116     df = df[df['trip_duration']>0]
117     df = df[df['tip_amt']>=0]
118
119     # One hot encoding of variables vendor_name and payment_type
120     cat_data = pd.get_dummies(df[['vendor_name', 'payment_type']])
121
122     # Target
123     y = df['fare_amt']
124
125     # Drop variables with no interest
126     df = df.drop(['index', 'trip_pickup_datetime',
127                 'trip_dropoff_datetime',
128                 'rate_code', 'mta_tax',
129                 'store_and_forward',
130                 'vendor_name', 'payment_type'], axis = 1)
131
132     # Concatenate dataframes
133     df = pd.concat([df, cat_data], axis=1)
134
135     return df
136
137 def pre_processing_2(df, scaler, numerical_variables):
138     # Min-max scaling of numerical variables
139     # numerical_variables=['passenger_count', 'trip_distance', 'start_lon', 'start_lat', '
140     # end_lon',
141     #                       'end_lat', 'surcharge', 'tip_amt', 'tolls_amt', 'total_amt',
142     #                       'trip_duration', 'hour']
143     #
144     #     scaler = MinMaxScaler()
145     #     scaler.fit(numerical_variables)
146
147     # Concatenate numerical variables and categorical variables
148     df = pd.concat([pd.DataFrame(scaler.transform(df[numerical_variables]), columns=
149     numerical_variables),
150     df.drop(numerical_variables, axis=1).reset_index(drop=True)], axis=1)
151
152     return df
153
154 # Training
155 from sklearn.model_selection import train_test_split
156 from xgboost import XGBRegressor
157
158 def training(X, y):
159     # Hold out of the data
160     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=0)
161
162     # Model training
163     model = XGBRegressor(objective='reg:squarederror')
164
165     # Model training
166     model.fit(X_train, y_train)
167
168     return model, X_train, y_train, X_test, y_test

```

```

168
169
170 def save_model(model, model_name):
171     io.save_obj(model, model_name)
172
173 # Prediction and test
174
175 def pred_test(model, X_test, y_test, mean_test):
176     # Predictions
177     prediction = model.predict(X_test)
178
179     # Evaluation metrics computation
180     df_error = y_test - prediction
181
182     # Metrics
183     df_metrics = pd.DataFrame()
184     df_metrics.loc[0, 'max_error'] = max(df_error)
185     df_metrics.loc[0, 'R^2'] = sum((y_test - mean_test)**2)
186     df_metrics.loc[0, 'MAE'] = sum(abs(df_error))
187     df_metrics.loc[0, 'MSE'] = sum(df_error**2)
188     df_metrics.loc[0, 'size'] = df_error.shape[0]
189
190     return df_metrics
191
192
193 def pipeline_joblib(n_jobs=-1):
194     t0 = time.time()
195     # Read data
196     print('[INFO] Reading data...')
197     df = read_data()
198     print('[INFO] Data loaded.')
199
200     # Pre-processing 1 (parallel)
201     tp0 = time.time()
202     print('[INFO] Pre-processing 1')
203     n_jobs = n_jobs if n_jobs > 0 else os.cpu_count()
204     dim0 = df.shape[0] # numero de linhas
205     chunk = math.ceil(dim0/n_jobs) #if n_jobs else math.ceil(dim0/os.cpu_count()) # tamanho do
206     chunk
207     nb_chunks = math.ceil(dim0/chunk)
208     index_list = [(i*chunk, min(chunk*(i+1), dim0)) for i in range(nb_chunks)]
209
210     res_list = Parallel(n_jobs=n_jobs)(delayed(pre_processing_1)(df.iloc[index[0]:index[1],:])
211     for index in index_list)
212     df = pd.concat(res_list)
213     # X = pd.concat([r[0] for r in res_list])
214     # y = pd.concat([r[1] for r in res_list])
215
216     tp1 = time.time()
217     print(f'[INFO] End of pre-processing 1 in {tp1-tp0}s.')
218
219     tp0 = time.time()
220     print(f'[INFO] Fit of the scaler.')
221     # Reset index (serial)
222     df = df.reset_index(drop=True)
223     # y = y.reset_index(drop=True)
224
225     # Fit of the Min-Max scaler (serial)
226     numerical_variables=['passenger_count', 'trip_distance', 'start_lon', 'start_lat', '
227     end_lon',
228     'end_lat', 'surcharge', 'tip_amt', 'tolls_amt', '#total_amt',
229     'trip_duration', 'hour', 'fare_amt']
230     scaler = MinMaxScaler()
231     scaler.fit(df[numerical_variables])
232
233     tp1 = time.time()
234     print(f'[INFO] End of the fit of the scaler in {tp1-tp0}s.')
235
236     # Pre-processing 2 (parallel)
237     tp0=time.time()
238     print('[INFO] Pre-processing 2')
239     dim0 = df.shape[0] # numero de linhas
240     chunk = math.ceil(dim0/n_jobs) #if n_jobs else math.ceil(dim0/os.cpu_count()) # tamanho do
241     chunk
242     nb_chunks = math.ceil(dim0/chunk)
243     index_list = [(i*chunk, min(chunk*(i+1), dim0)) for i in range(nb_chunks)]

```

```

240 res_list = Parallel(n_jobs=n_jobs)(delayed(pre_processing_2)(df.iloc[index[0]:index[1],:],
241     scaler=scaler, numerical_variables=numerical_variables) for index in index_list)
242 df = pd.concat(res_list)
243 tp1 = time.time()
244 print(f'[INFO] End of pre-processing 2 in {tp1-tp0}s.')
245
246 # Training (serial function, but XGBR training is parallel)
247 print('[INFO] Training the model...')
248 tp0 = time.time()
249 model, X_train, y_train, X_test, y_test = training(df.drop(['fare_amt'],axis=1), df['
250 fare_amt'])
251 tp1 = time.time()
252 print(f'[INFO] Model trained in {tp1-tp0} s.')
253
254 # Save the model (serial)
255 print('[INFO] Saving the model...')
256 tp0 = time.time()
257 save_model(model, folder_out + 'Joblib_model'+ datetime.now().strftime("%H%M%S"))
258 tp1 = time.time()
259 print('[INFO] Model saved.')
260
261 # Predict and model metrics (parallel)
262 tp0 = time.time()
263 print('[INFO] Predicting and computing model metrics...')
264 mean_test = func_parallel(mean, n_jobs=-1, df=y_test)
265
266 dim0 = X_test.shape[0] # numero de linhas
267 chunk = math.ceil(dim0/n_jobs) #if n_jobs else math.ceil(dim0/os.cpu_count()) # tamanho do
268 chunk
269 nb_chunks = math.ceil(dim0/chunk)
270 index_list = [(i*chunk,min(chunk*(i+1),dim0)) for i in range(nb_chunks)]
271
272 res_list = Parallel(n_jobs=n_jobs)(delayed(pred_test)(model,
273     X_test.iloc[index[0]:index[1],:],
274     y_test.iloc[index[0]:index[1]],
275     mean_test) for index in index_list)
276 res = functools.reduce(lambda a, b: a.add(b, fill_value=0), res_list)
277 res = res/res['size'].values[0]
278 res['max_error'] = max([x['max_error'].values[0] for x in res_list])
279 res['R^2'] = 1 - sum([x['MSE'].values[0] for x in res_list])/sum([x['R^2'].values[0] for x
280 in res_list])
281 res.loc[0,'RMSE'] = (res.loc[0,'MSE'])**(1/2)
282
283 res = res.drop(['size'],axis=1)
284
285 # Save metrics
286 io.save_obj(res,folder_out + 'Joblib_model_metrics'+ datetime.now().strftime("%H%M%S"))
287 tp1 = time.time()
288 print(f'[INFO] Metrics computed and saved in {tp1-tp0} s.')
289
290 t1 = time.time()
291 print(f'[INFO] All metrics computed.\n[INFO] End of the pipeline in {t1-t0} seconds.')

```

E.3 Koalas

```

1
2 import pandas as pd
3 import numpy as np
4 import databricks.koalas as ks
5 from datetime import timedelta, datetime
6 from utils import io
7 import time
8
9 import matplotlib.pyplot as plt
10
11 from pyspark.ml.regression import LinearRegression
12 from pyspark.ml.feature import VectorAssembler
13 from pyspark.ml.feature import MinMaxScaler
14 from pyspark.ml.evaluation import RegressionEvaluator
15 from pyspark.ml.regression import GBRegressor
16 from pyspark.ml.linalg import Vectors
17 from pyspark.sql.functions import udf
18 from pyspark.sql.types import DoubleType
19 from pyspark.sql.functions import monotonically_increasing_id

```

```

20
21 #####
22 ##### Read data #####
23 #####
24 def read_data():
25     koalas_data = ks.read_parquet('.././data/2009-01', index_col='index')
26     return koalas_data
27
28
29 #####
30 ##### Pre-processing #####
31 #####
32 def pre_processing(koalas_data):
33
34     koalas_data=koalas_data.assign(pdayofweek = koalas_data.trip_pickup_datetime.dt.dayofweek,
35                                   phour = koalas_data.trip_pickup_datetime.dt.hour)
36
37     # Creation of the variable weekend (0-weekend, 1-work day)
38     koalas_data['WEEKDAY'] = koalas_data['pdayofweek'] //5
39
40     # Drop of variable weekend
41     koalas_data = koalas_data.drop(['pdayofweek'], axis = 1)
42
43     # Creation of the variables trip_duration
44     koalas_data['trip_duration'] = koalas_data['trip_dropoff_datetime'] - koalas_data['trip_pickup_datetime']
45
46     # Recode same variables with different names
47     koalas_data['payment_type'] = koalas_data['payment_type'].mask(koalas_data['payment_type']
48                             == 'CASH', 'Cash')
49     koalas_data['payment_type'] = koalas_data['payment_type'].mask(koalas_data['payment_type']
50                             == 'CREDIT', 'Credit')
51
52     # Remove lines with abnormal values
53     koalas_data = koalas_data[koalas_data['fare_amt']>0]
54     koalas_data = koalas_data[koalas_data['trip_distance']>0]
55     koalas_data = koalas_data[koalas_data['trip_duration']>0]
56     koalas_data = koalas_data[koalas_data['tip_amt']>=0]
57
58     # One hot encoding of variables vendor_name and payment_type
59     cat_data = ks.get_dummies(koalas_data[['vendor_name', 'payment_type']])
60
61     # Transform Categorical Variables in Koalas Dataframe to Pyspark Dataframe
62     pycat = cat_data.to_spark()
63
64     # Drop variables with no interest
65     num_data = koalas_data.drop(['index', 'trip_pickup_datetime', 'trip_dropoff_datetime',
66                                 'rate_code', 'mta_tax', 'store_and_forward', 'total_amt',
67                                 'vendor_name', 'payment_type'], axis = 1)
68
69     # Transform Numerical Variables Koalas Dataframe to Pyspark Dataframe
70     pynum = num_data.to_spark()
71
72     # Adding all values except "fare_amt" to an array called "features"
73     vectorAssembler_num = VectorAssembler(inputCols=['passenger_count', 'trip_distance', 'start_lon', 'start_lat',
74                                                       'end_lon', 'end_lat', 'surcharge', 'tip_amt',
75                                                       'tolls_amt',
76                                                       'phour', 'trip_duration', 'WEEKDAY'],
77                                           outputCol = "num_features")
78     pynum2 = vectorAssembler_num.transform(pynum)
79
80     # Apply MinMaxScaler to num_features and return scaledFeatures column
81     scaler = MinMaxScaler(inputCol="num_features", outputCol="scaledFeatures")
82     scalerModel = scaler.fit(pynum2)
83     scaledData = scalerModel.transform(pynum2)
84
85     # Transform fare_amt to a vector to apply MinMaxScaler on target variable
86     vectorAssembler_fare = VectorAssembler(inputCols=['fare_amt'], outputCol = "fare_amt_f")
87     scaledData2 = vectorAssembler_fare.transform(scaledData)
88
89     # Apply MinMaxScaler to fare_amt_f and return s_fare_amt column
90     scaler2 = MinMaxScaler(inputCol="fare_amt_f", outputCol="s_fare_amt")
91     scalerModel2 = scaler2.fit(scaledData2)
92     scaledData3 = scalerModel2.transform(scaledData2)

```

```

91 # Create id column on both categorical and numerical dataframes and merge by id
92 df1 = scaledData3
93 df2 = pycat
94
95 df1 = df1.withColumn("id", monotonically_increasing_id())
96 df2 = df2.withColumn("id", monotonically_increasing_id())
97
98 df3 = df2.join(df1, "id").drop("id")
99
100 # Add categorical values and scaledFeatures into new column features
101 vectorAssembler = VectorAssembler(inputCols = ['vendor_name_CMT', 'vendor_name_DDS', '
102                                     'payment_type_Cash', 'payment_type_Credit',
103                                     'payment_type_Dispute', 'payment_type_No Charge'
104                                     ,
105                                     'scaledFeatures'], outputCol = "
106 features")
107 df4 = vectorAssembler.transform(df3)
108
109 # Convert vector s_fare_amt to Double scaled_fare_amt
110 unlist = udf(lambda x: float(list(x)[0]), DoubleType())
111 pyko2 = df4.withColumn("scaled_fare_amt", unlist("s_fare_amt"))
112
113 return koalas_data, pyko2
114
115 #####
116 ##### Training #####
117 #####
118 def train(pyko2):
119     # Hold out of the data
120     splits = pyko2.randomSplit([0.7,0.3])
121     train_df = splits[0]
122     test_df = splits[1]
123
124     # Model training
125     gbt = GBTRegressor(featuresCol="features", labelCol="scaled_fare_amt", lossType = "squared
126 ")
127
128     # Model training
129     gbt_model = gbt.fit(train_df)
130
131     return gbt_model, train_df, test_df
132
133 #####
134 ##### Prediction and test #####
135 #####
136
137 def pred_test(gbt_model, test_df):
138     # Predictions
139     gbt_predictions = gbt_model.transform(test_df)
140
141     # Root Mean Square Error
142     gbt_evaluator_rmse = RegressionEvaluator(labelCol="scaled_fare_amt", predictionCol="
143 prediction", metricName="rmse")
144     gbt_rmse = gbt_evaluator_rmse.evaluate(gbt_predictions)
145
146     # Mean Square Error
147     gbt_evaluator_mse = RegressionEvaluator(labelCol="scaled_fare_amt", predictionCol="
148 prediction", metricName="mse")
149     gbt_mse = gbt_evaluator_mse.evaluate(gbt_predictions)
150
151     # Mean Absolute Error
152     gbt_evaluator_mae = RegressionEvaluator(labelCol="scaled_fare_amt", predictionCol="
153 prediction", metricName="mae")
154     gbt_mae = gbt_evaluator_mae.evaluate(gbt_predictions)
155
156     # R2
157     gbt_evaluator_r2 = RegressionEvaluator(labelCol="scaled_fare_amt", predictionCol="
158 prediction", metricName="r2")
159     gbt_r2 = gbt_evaluator_r2.evaluate(gbt_predictions)
160
161     # Kolas Dataframe with predictions
162     koalas_predict = ks.DataFrame(gbt_predictions)

```



```

159
160 # Metrics
161 df_metrics = pd.DataFrame()
162 df_metrics.loc[0,'max_error'] = ((koalas_predict['prediction']-koalas_predict['
scaled_fare_amt']).abs()).max()
163 df_metrics.loc[0,'R^2'] = gbt_r2
164 df_metrics.loc[0,'MAE'] = gbt_mae
165 df_metrics.loc[0,'MSE'] = gbt_mse
166 df_metrics.loc[0,'RMSE'] = gbt_rmse
167
168 return koalas_predict, df_metrics
169
170 #####
171 ##### Pipeline #####
172 #####
173
174 def pipeline_koalas():
175     t0 = time.time()
176
177     # Read data
178     tr0 = time.time()
179     print('[INFO] Reading data...')
180     koalas_data = read_data()
181     tr1 = time.time()
182     print(f'[INFO] Data loaded in {tr1-tr0} s.\n')
183
184     # Pre processing
185     tp0 = time.time()
186     print('[INFO] Pre-processing')
187     koalas_data, pyko2 = pre_processing(koalas_data)
188     tp1 = time.time()
189     print(f'[INFO] Pre-processing done in {tp1-tp0} s.\n')
190
191     # Training
192     tt0 = time.time()
193     print('[INFO] Training the model...')
194     gbt_model, train_df, test_df = train(pyko2)
195     tt1 = time.time()
196     print(f'[INFO] Model trained in {tt1-tt0} s.\n')
197
198     # Predict and model metrics
199     tm0 = time.time()
200     print('[INFO] Predicting and computing model metrics...')
201     koalas_predict, df_metrics = pred_test(gbt_model, test_df)
202     tm1 = time.time()
203     print(f'[INFO] Predictions and all metrics computed in {tm1-tm0} s.\n')
204
205     # Save metrics
206     io.save_obj(df_metrics, 'Koalas_model_metrics'+ datetime.now().strftime("%H%M%S"))
207
208     t1 = time.time()
209     print(f'[INFO] Pipeline computed in {t1-t0} s.\n[INFO] End of the pipeline.')
210     return koalas_predict, df_metrics
211
212 #####
213 ##### Yappi #####
214 #####
215
216 import yappi
217
218 # Yappi Read
219 yappi.clear_stats()
220 yappi.start()
221 koalas_data = read_data()
222 yappi.stop()
223 stats = yappi.get_func_stats(filter_callback=lambda x: x.ttot>0.5)#.print_all()
224 stats.save('../outputs/callgrind.kread.prof', type = "ystat")
225
226 # Yappi Pre Processing
227 yappi.clear_stats()
228 yappi.start()
229 koalas_data, pyko2 = pre_processing(koalas_data)
230 yappi.stop()
231 stats = yappi.get_func_stats(filter_callback=lambda x: x.ttot>0.5)#.print_all()
232 stats.save('../outputs/callgrind.kprepros.prof', type = "ystat")
233

```

```

234 # Yappi Train
235 yappi.clear_stats()
236 yappi.start()
237 gbt_model, train_df, test_df = train(pyko2)
238 yappi.stop()
239 stats = yappi.get_func_stats(filter_callback=lambda x: x.ttot>0.5)#.print_all()
240 stats.save('../outputs/callgrind.ktrain.prof', type = "ystat")
241
242 # Yappi Pred
243 yappi.clear_stats()
244 yappi.start()
245 koalas_predict, df_metrics = pred_test(gbt_model, test_df)
246 yappi.stop()
247 stats = yappi.get_func_stats(filter_callback=lambda x: x.ttot>0.5)#.print_all()
248 stats.save('../outputs/callgrind.kpred.prof', type = "ystat")

```

E.4 Rapids

```

1 def import_data(path = None):#"../data/ks_taxi_parquet"
2     t0 = time.time()
3
4     try:
5         rapids_data = cudf.io.parquet.read_parquet(path)
6         global train_ind
7         train_ind = round(len(rapids_data)*0.7)
8     except:
9         print("No path provided.")
10    t1 = time.time()
11
12    print(f'[INFO] Computation time for loading the file: {t1-t0}s')
13    return rapids_data
14
15
16 def preprocessing(rapids_data):
17     t0 = time.time()
18
19     rapids_data = rapids_data.assign(pdayofweek = rapids_data.trip_pickup_datetime.dt.
20                                     dayofweek,
21                                     phour = rapids_data.trip_pickup_datetime.dt.hour)
22
23     # computation of the variable weekend
24     #zero weekend, 1 week
25     rapids_data['weekend'] = (rapids_data['pdayofweek'] < 5).astype(int)
26
27     # computation of the variable trip_duration
28     rapids_data['trip_duration'] = rapids_data['trip_dropoff_datetime'] - rapids_data['
29     trip_pickup_datetime']
30     rapids_data['trip_duration'] = rapids_data['trip_duration']/timedelta(seconds=1)
31
32     # Corretion of the classes of the variable payment_type
33     rapids_data['payment_type'] = rapids_data['payment_type'].mask(rapids_data['payment_type'
34     ]=='CASH', 'Cash')
35     rapids_data['payment_type'] = rapids_data['payment_type'].mask(rapids_data['payment_type'
36     ]=='CREDIT', 'Credit')
37
38     # Drop concatentated timestamp columns
39     rapids_data = rapids_data.drop(['trip_pickup_datetime', 'trip_dropoff_datetime', '
40     rate_code', 'mta_tax', 'store_and_forward', 'pdayofweek'], axis = 1)
41
42     # Remove lines with abnormal values
43     rapids_data = rapids_data[rapids_data['fare_amt']>0]
44     rapids_data = rapids_data[rapids_data['trip_distance']>0]
45     rapids_data = rapids_data[rapids_data['trip_duration']>0]
46     rapids_data = rapids_data[rapids_data['tip_amt']>=0]
47     rapids_data = rapids_data.drop('total_amt', axis = 1)
48
49     rapids_data = rapids_data.iloc[:,1:].reset_index()
50
51     # One hot encoding of variables vendor_name and payment_type
52     rapids_data["vendor_name"] = rapids_data["vendor_name"].astype("category")
53     rapids_data["payment_type"] = rapids_data["payment_type"].astype("category")
54
55     categorical_variables = rapids_data[['vendor_name', 'payment_type']]
56
57     # encoder = OneHotEncoder(sparse=False)

```

```

53 # encoder.fit(categorical_variables)
54 # cat_data = encoder.transform(categorical_variables).compute()
55
56 cat_data = cudf.get_dummies(categorical_variables)
57
58 # Min-max scaler of numerical variables
59 numeric_variables = rapids_data[['passenger_count', 'trip_distance', 'start_lon', '
60 start_lat', 'end_lon',
61                                     'end_lat', 'surcharge', 'tip_amt', 'tolls_amt', '
62 trip_duration', 'fare_amt']]
63 scaler = MinMaxScaler()
64 scaler.fit(numeric_variables)
65 num_data = scaler.transform(numeric_variables)
66
67 num_data = num_data.rename(columns={0: 'passenger_count', 1: 'trip_distance', 2: '
68 start_lon', 3: 'start_lat', 4: 'end_lon', 5: 'end_lat',
69                                     6: 'surcharge', 7: 'tip_amt', 8: 'tolls_amt', 9: '
70 trip_duration', 10: 'fare_amt'})
71
72 rapids_data = cudf.concat([num_data, cat_data], axis=1) #rapids_data['fare_amt']
73
74 t1 = time.time()
75 print(f'[INFO] Computation time for preprocessing the data: {t1-t0}s')
76 return rapids_data
77
78 def training(X, y):#X = rapids_data.drop('fare_amt', axis = 1), y = rapids_data['fare_amt']
79 # Hold out of the data
80 X_train = X[0:train_ind]
81 X_test = X[train_ind:]
82 y_train = y[0:train_ind]
83 y_test = y[train_ind:]
84
85 #Conversion to a DMatrix
86 X_train_DM = xgb.DMatrix(X_train.values, label=y_train.values)
87
88 # Parameters and Model training
89 param = {'objective': 'reg:squarederror'}
90
91 t0 = time.time()
92
93 bst = xgb.train(param, X_train_DM)
94
95 t1 = time.time()
96
97 print(f'[INFO] Computation time for training the model: {t1-t0}s')
98 return bst, X_train, y_train, X_test, y_test
99
100 from datetime import datetime
101 def save_model(model, model_name):
102     io.save_obj(model, model_name + datetime.now().strftime("%H%M%S"))
103
104 import cuml.metrics.regression as cuml_metrics
105 import sklearn.metrics as skl_metrics
106 # Prediction and test
107
108 def pred_test(model, X_test, y_test):
109     #Conversion to a DMatrix
110     X_test_DM = xgb.DMatrix(X_test.values, label = y_test.values)
111
112     # Predictions
113     prediction = model.predict(X_test_DM)
114
115     # Evaluation metrics computation
116     max_error = skl_metrics.max_error(y_test.to_array(), prediction)
117     variance = cuml_metrics.r2_score(y_test.to_array(), prediction)
118     mean_absolute_error = cuml_metrics.mean_absolute_error(np.float64(y_test.to_array()), np.
119 float64(prediction))
120     mean_squared_error = cuml_metrics.mean_squared_error(np.float64(y_test.to_array()), np.
121 float64(prediction))
122     root_mean_squared = np.sqrt(cuml_metrics.mean_squared_error(np.float64(y_test.to_array()),
123 np.float64(prediction)))
124
125 print('Max Error:', max_error)

```

```

122 print('R^2:', variance)
123 print('MAE:', mean_absolute_error)
124 print('MSE:', mean_squared_error)
125 print('RMSE:', root_mean_squared)
126
127 # Metrics
128 gdf_metrics = cudf.DataFrame({'Max Error' : max_error,
129                               'Explained Variance': variance,
130                               'Mean absolute error': mean_absolute_error,
131                               'Mean squared error': mean_squared_error,
132                               'Root mean Squared error': root_mean_squared})
133
134 io.save_obj(gdf_metrics, '../outputs/Metrics/Rapids_model_metrics'+ datetime.now().
135             strftime("%H%M%S"))
136 return gdf_metrics
137
138 def pipeline_rapids():
139     t0 = time.time()
140     # Read data
141     df = import_data("../data/2009-01")
142
143     df = preprocessing(df)
144     # Training (serial function, but XGBR training is parallel)
145     model, X_train, y_train, X_test, y_test = training(df.drop('fare_amt', axis = 1), df['
146     fare_amt'])
147
148     # Save the model (serial)
149     save_model(model, '../outputs/Models/Rapids_model')
150
151     # Predict and model metrics
152     df_metrics = pred_test(model, X_test, y_test)
153
154     t1 = time.time()
155     print(f'[INFO] Computation time to perform the pipeline: {t1-t0}s')
156     return df_metrics
157
158 pipeline_rapids()
159
160 %load_ext line_profiler
161
162 %lprun -f pipeline_rapids pipeline_rapids()
163
164 import yappi
165
166 yappi.clear_stats()
167
168 yappi.start()
169 import_data("../data/2009-01")
170 yappi.stop()
171
172 stats_import_data = yappi.get_func_stats(filter_callback=lambda x: x.ttot>2.5)#.print_all()
173 stats_import_data.save('../outputs/Yappi/ystat.import_data' + datetime.now().strftime("%H%M%S"
174 ) + '.prof', type = "ystat")
175
176 #stats.save('../outputs/Dask/callgrind.preproc.' + datetime.now().isoformat(), 'callgrind')
177
178 import yappi
179
180 df = import_data("../data/2009-01")
181
182 yappi.clear_stats()
183
184 yappi.start()
185 df = preprocessing(df)
186 yappi.stop()
187
188 stats_preprocessing = yappi.get_func_stats(filter_callback=lambda x: x.ttot>2.5)#.print_all()
189 stats_preprocessing.save('../outputs/Yappi/ystat.preprocessing' + datetime.now().strftime("%H%
190 M%S") + '.prof', type = "ystat")
191
192 #stats.save('../outputs/Dask/callgrind.preproc.' + datetime.now().isoformat(), 'callgrind')
193

```

```

194 import yappi
195
196 yappi.clear_stats()
197
198 yappi.start()
199 model, X_train, y_train, X_test, y_test = training(df.drop('fare_amt', axis = 1), df['fare_amt
200 yappi.stop()
201
202 stats_training = yappi.get_func_stats(filter_callback=lambda x: x.ttot>2.5)#.print_all()
203 stats_training.save('../outputs/Yappi/ystat.training' + datetime.now().strftime("%H%M%S") + '.
    prof', type = "ystat")
204
205 #stats.save('../outputs/Dask/callgrind.preproc.' + datetime.now().isoformat(), 'callgrind')
206
207
208 import yappi
209
210 yappi.clear_stats()
211
212 yappi.start()
213 save_model(model, 'Rapids_model_yappi')
214 yappi.stop()
215
216 stats_save_model = yappi.get_func_stats(filter_callback=lambda x: x.ttot>0.1)#.print_all()
217 stats_save_model.save('../outputs/Yappi/ystat.save_model' + datetime.now().strftime("%H%M%S")
    + '.prof', type = "ystat")
218
219 #stats.save('../outputs/Dask/callgrind.preproc.' + datetime.now().isoformat(), 'callgrind')
220
221
222 import yappi
223
224 yappi.clear_stats()
225
226 yappi.start()
227 pred_test(model, X_test, y_test)
228 yappi.stop()
229
230 stats_prediction = yappi.get_func_stats(filter_callback=lambda x: x.ttot>0.5)#.print_all()
231 stats_prediction.save('../outputs/Yappi/ystat.prediction' + datetime.now().strftime("%H%M%S")
    + '.prof', type = "ystat")
232
233 #stats.save('../outputs/Dask/callgrind.preproc.' + datetime.now().isoformat(), 'callgrind')
234
235
236 import yappi
237
238 yappi.clear_stats()
239
240 yappi.start()
241 pipeline_rapids()
242 yappi.stop()
243
244 stats_pipeline = yappi.get_func_stats(filter_callback=lambda x: x.ttot>10)#.print_all()
245 stats_pipeline.save('../outputs/Yappi/ystat.pipeline' + datetime.now().strftime("%H%M%S") + '.
    prof', type = "ystat")
246
247 #stats.save('../outputs/Dask/callgrind.preproc.' + datetime.now().isoformat(), 'callgrind')

```